
Finite-Support Convolutional Likelihoods for Aggregate-Supervised Multiple-Instance Learning

Anonymous Author(s)

Affiliation

Address

email

Abstract

Weakly supervised multiple-instance learning often observes only aggregate labels for a bag, while instance labels remain hidden. Existing count-based objectives give exact likelihoods for Bernoulli counts, but many aggregate labels have richer finite structure, such as signed counts, integer-valued sums, and class histograms. We formulate aggregate-supervised Count-MIL as a finite-support convolutional aggregator: an architecture-independent layer that maps instance-level atomic probability mass functions to the mathematically implied aggregate distribution. The resulting bag likelihood recovers the standard Bernoulli count probability as a special case and extends, without changing the backbone model or the aggregate layer, to signed and integer-valued aggregate labels. We implement exact aggregate inference with GPU-parallel convolution, including an FFT-tree backend for large one-vs-rest histogram likelihoods. Experiments on digit sums, signed counts, Bernoulli counts, SVHN/UltraMNIST sum benchmarks, and class-count settings show that exact finite-support likelihoods can recover accurate hidden-instance predictors from aggregate labels. On CIFAR-10, ImageNet-pretrained ResNet-18 features make the one-vs-rest count likelihood effective, while a fixed-instance-budget comparison shows that the large-bag improvement is explained by greater sampled image exposure.

1 Introduction

Multiple-instance learning (MIL) studies settings where supervision is observed for a bag while the labels of individual instances remain latent. Classical binary MIL uses a positive bag label to indicate that at least one instance is positive. Many weak-supervision settings provide more structured aggregate labels: the number of positive instances, a class histogram, a signed count, or a scalar integer-valued sum. The central problem is to learn an instance-level predictor from such aggregate observations without using instance labels during training.

This paper studies aggregate-supervised MIL through a finite-support convolutional aggregator. Each instance x_i is mapped to a probability mass function over a finite set of integer-valued contributions. The aggregator then returns the mathematically implied distribution of the observed bag aggregate by convolving the instance-level atomic PMFs. For Bernoulli contributions this recovers the standard count probability used in prior count-loss work. The same construction also gives exact likelihoods for signed counts, digit-sum labels, and one-vs-rest class-count objectives. We refer to this as finite-support aggregation: the atomic support size and contribution values define the weak-label problem, while the convolutional aggregator is unchanged. The formulation separates the instance model from the aggregate likelihood: a backbone model outputs instance-level scores, and the aggregate layer converts these scores into finite-support atomic PMFs before computing the bag likelihood. The method is therefore not tied to a particular architecture. Changing the weak




Figure 1: Convolutional aggregate pipeline. The instance network emits finite-support atomic PMFs d_i for each instance. A deterministic convolutional layer computes the aggregate PMF $P_\theta(Y)$, and the observed bag label selects the likelihood term used for training.

label changes the contribution map and convolution support, not the backbone, the optimizer, or the probabilistic rule used to obtain the bag-level distribution.

This view also clarifies the boundary with prior work. Binary Bernoulli counts recover existing count-loss objectives, and one-vs-rest class-count likelihoods are closely related to the count-likelihood component of LLP-PVC. Our focus is the finite-support aggregate layer that covers these cases while also extending to signed and integer-valued sum labels.

Our contributions are:

- We formulate aggregate-supervised Count-MIL as a backbone-agnostic finite-support convolutional aggregator that turns instance contribution PMFs into an exact aggregate PMF, explicitly recovering Bernoulli count loss as a special case rather than claiming it as new.
- We show that the same backbone-agnostic likelihood layer extends to signed counts, digit-sum labels, and one-vs-rest histogram class counts by changing only the finite-support contribution map.
- We implement practical convolutional backends, including grouped signed convolution and a GPU FFT-tree backend that parallelizes large multiclass one-vs-rest PMFs beyond CPU dynamic programming.
- We evaluate hidden-instance recovery across binary, signed, digit-sum, SVHN/UltraMNIST sum, and one-vs-rest class-count settings, using histogram experiments mainly to connect the formulation to LLP-PVC-style count likelihoods.

Empirically, the largest improvements occur for aggregate types where expectation matching discards information, especially digit sums and signed counts. The CIFAR-10 comparison shows that count likelihood achieves high instance accuracy when paired with ImageNet-pretrained features, but the fixed-instance-budget result attributes the large-bag improvement to greater sampled image exposure rather than to bag size alone.

2 Related Work

Multiple-instance learning and count loss. Classical multiple-instance learning supervises bags rather than instances [8, 16, 1, 6]; neural and attention MIL learn bag predictors end-to-end [23, 11]. Count-based weak supervision gives a stronger bag label by observing the number of positive instances. Shukla et al. [20] compute the exact probability that k Bernoulli outputs are true, and related count supervision appears in temporal localization [18]. Our Bernoulli case recovers this count probability; the contribution is the finite-support convolutional formulation that also covers signed, integer-valued, and classwise histogram aggregates.

Learning from label proportions. Learning from label proportions observes class proportions or histograms over bags [17, 26], with probabilistic, deep, pseudo-labeling, consistency, and two-stage

variants [21, 3, 10, 9, 12, 22, 13, 25, 14] and recent theoretical analyses [19, 27, 4, 24, 2, 5]. LLP-PVC [15] reframes label proportions as proportional-value classification and uses efficient count-likelihood computation. Our one-vs-rest histogram experiment is included to connect to this count-likelihood component, not to claim a new full joint histogram likelihood; the broader contribution is the same aggregate layer applied to scalar sums and signed counts.

Efficient aggregate inference. Dynamic programming computes Bernoulli count probabilities exactly but is sequential and can be expensive for large bags and many classes [15]. The same finite-support likelihood can instead be evaluated by direct convolution or FFT polynomial multiplication [7]. We use sequential convolution for moderate aggregate supports and a GPU FFT-tree backend for large one-vs-rest count PMFs; this changes the backend, not the statistical objective.

3 Problem Setup

We study weakly supervised bags of instances. A training example is a bag $B = (x_1, \dots, x_n)$ together with an observed bag label y . The instance-level variables that explain y are hidden during training and are used only for evaluation. The label y may be a binary count, a signed count, an integer-valued sum, or a class histogram. The common structure is that each instance contributes a finite-support discrete value and the bag label is a sum of these contributions. Our goal is to define the aggregate layer once, so that different weak-supervision problems are obtained by changing only the atomic support and contribution map.

The unifying object is the *atomic support*. For a one-dimensional task, choose a finite set

$$\mathcal{A} = \{v_0, \dots, v_{S-1}\} \subset \mathbb{Z}, \quad (1)$$

where $S = |\mathcal{A}|$ is the atomic support size. A backbone model produces scores for each instance, and a contribution map turns those scores into an atomic PMF over $\mathcal{A}$. Changing $\mathcal{A}$ changes the weak-label problem while leaving the bag-prediction algorithm unchanged: $S = 2$ gives Bernoulli counts, $S = 10$ gives digit sums over $\{0, \dots, 9\}$, and $S = 3$ gives signed atoms over $\{-1, 0, +1\}$. For bag size n and contiguous support $\{0, \dots, S-1\}$, the aggregate PMF has length

$$T = n(S-1) + 1. \quad (2)$$

Signed supports use the same principle with an offset, and variable-multiplicity atoms use instance-dependent support sizes.

Let C_i denote the latent contribution of instance x_i . The contribution may also depend on observed side information a_i , such as a known sign in signed-count experiments. We assume

$$C_i \in \mathcal{S}_i \subset \mathbb{Z}^d, \quad |\mathcal{S}_i| < \infty, \quad (3)$$

and the observed aggregate is

$$Y = \sum_{i=1}^n C_i. \quad (4)$$

For most experiments $d = 1$. For histograms, d can be the number of classes, although our scalable implementation uses one-dimensional one-vs-rest count likelihoods. The learner observes $(x_{1:n}, a_{1:n}, y)$ when side information is present, but not the individual C_i values. When discussing individual scalar tasks, we write Y_{count} for ordinary nonnegative counts, Y_{sum} for integer-valued sums such as digit sums, and $Y_{\pm}$ for signed counts. These are task-specific instances of the generic aggregate Y in Eq. (4).

Definition 1 (Finite-support convolutional aggregator). *Let f_θ be any instance model, and let Φ be a contribution map that converts the output of f_θ on instance x_i and observed side information a_i into an atomic PMF*

$$d_i = \Phi(f_\theta(x_i), a_i) \in \Delta(\mathcal{S}_i), \quad \mathcal{S}_i \subset \mathbb{Z}^d, \quad |\mathcal{S}_i| < \infty. \quad (5)$$

The finite-support convolutional aggregator is the deterministic map

$$\text{Agg}(d_1, \dots, d_n)(y) = (d_1 * \dots * d_n)(y), \quad (6)$$

which assigns probability mass to aggregate value y .

111 Definition 1 separates four choices often coupled in task-specific weak-supervision losses:

- 112 • *Backbone independence.* The network supplies atomic PMFs; the aggregate layer is un-
113 changed across MLPs, CNNs, transformers, or pretrained encoders.
- 114 • *Kernel independence.* Binary counts, signed counts, digit sums, and one-vs-rest class
115 counts are specified by Φ and $\mathcal{S}_i$, not by new bag models.
- 116 • *Objective and backend independence.* The aggregator returns a full PMF that can be trained
117 with NLL and evaluated by DP, direct convolution, grouped convolution, or FFT-tree mul-
118 tiplication.
- 119 • *Conditional-independence scope.* The PMF is exact under conditionally independent latent
120 contributions; otherwise it is a structured aggregate-likelihood approximation.

121 In this paper we use the aggregate negative log likelihood $-\log \text{Agg}(d_1, \dots, d_n)(y)$ for the observed
122 weak label y .

123 **Proposition 1** (Finite-support convolution law). *Let $C_1, \dots, C_n$ be conditionally independent dis-
124 crete random variables given a bag $x_{1:n}$ and any observed side information $a_{1:n}$. Assume each C_i
125 has finite support $\mathcal{S}_i \subset \mathbb{Z}^d$ and probability mass function*

$$d_i(c) = \Pr_{\theta}(C_i = c \mid x_i, a_i), \quad c \in \mathbb{Z}^d, \quad (7)$$

126 where $d_i(c) = 0$ for $c \notin \mathcal{S}_i$. Define $Y = \sum_{i=1}^n C_i$. Then Y has finite support contained in the
127 Minkowski sum

$$\mathcal{S}_1 + \dots + \mathcal{S}_n = \{c_1 + \dots + c_n : c_i \in \mathcal{S}_i\}, \quad (8)$$

128 and its probability mass function is the discrete convolution

$$\Pr_{\theta}(Y = y \mid x_{1:n}, a_{1:n}) = (d_1 * \dots * d_n)(y) = \sum_{\substack{c_1 \in \mathcal{S}_1, \dots, c_n \in \mathcal{S}_n \\ c_1 + \dots + c_n = y}} \prod_{i=1}^n d_i(c_i). \quad (9)$$

129 Proposition 1 shows why Definition 1 is the mathematically correct aggregate PMF under the
130 conditional-independence model. The aggregator is not a heuristic pooling rule: once the atomic
131 contribution distributions d_i are predicted, the bag-label distribution is determined exactly by con-
132 volution. Different weak-label tasks differ mainly in the support $\mathcal{S}_i$, the contribution map Φ , and
133 which side information is observed; concrete kernels are summarized in Table 1. Binary Count-
134 MIL is therefore the simplest case of Eq. (4), while signed counts, integer-valued sums, multiplicity
135 counts, and one-vs-rest histograms change the finite-support kernel but keep the same aggregator. In
136 our experiments, bags are sampled on the fly with replacement from the corresponding train or test
137 split. Thus an image may appear in multiple bags, but train and test bags are always drawn from
138 disjoint dataset splits.

139 4 Convolutional Bag Prediction

140 4.1 Aggregate Likelihood

141 Following Definition 1, the instance network and contribution map emit the atomic PMFs d_i . For
142 one-dimensional contributions, this is a categorical distribution over a finite integer support:

$$d_i(k) = p_{\theta}(z_i = k \mid x_i), \quad k \in \{a, \dots, b\}. \quad (10)$$

143 By Proposition 1, the exact bag-label likelihood is the convolutional aggregator applied to these
144 atomic PMFs:

$$p_{\theta}(Y = y \mid x_{1:n}) = \text{Agg}(d_1, \dots, d_n)(y) = (d_1 * d_2 * \dots * d_n)(y). \quad (11)$$

145 Training minimizes the negative log likelihood

$$\mathcal{L}_{\text{agg}}(\theta) = -\log p_{\theta}(Y = y \mid x_{1:n}). \quad (12)$$

146 The support of the aggregate PMF is finite and contiguous:

$$\text{supp}(Y) = \{na, na + 1, \dots, nb\}. \quad (13)$$

147 Equivalently, if each atom has contiguous support size $S = b - a + 1$, then the aggregate PMF
148 has width $T = n(S - 1) + 1$ after shifting by na . All experiments use exact likelihoods over this
149 support; approximations enter only through the neural instance model, not through sampling latent
150 assignments.

Algorithm 1 Convolutional bag prediction for one minibatch

Require: Bags $\{(x_{b1:n_b}, a_{b1:n_b}, y_b)\}_{b=1}^B$, backbone f_θ , contribution map Φ , backend Conv

Ensure: Loss $\mathcal{L}$ and bag-label probabilities $\{p_b\}_{b=1}^B$

```
1: for  $b = 1, \dots, B$  do
2:    $H_b \leftarrow f_\theta(x_{b1:n_b})$  ▷ instance logits or class scores
3:   for  $i = 1, \dots, n_b$  do
4:      $d_{bi} \leftarrow \Phi(H_{bi}, a_{bi})$  ▷ atomic PMF on the task support
5:   end for
6:   if histogram one-vs-rest objective then
7:     for  $c = 1, \dots, C$  do
8:        $m_{bc} \leftarrow \text{Conv}(\{[1 - p_{bic}, p_{bic}]\}_{i=1}^{n_b})$ 
9:        $p_{bc} \leftarrow m_{bc}(h_{bc})$ 
10:    end for
11:     $p_b \leftarrow \prod_{c=1}^C p_{bc}^{1/C}$ 
12:     $\ell_b \leftarrow -C^{-1} \sum_{c=1}^C \log p_{bc}$ 
13:  else
14:     $m_b \leftarrow \text{Conv}(\{d_{bi}\}_{i=1}^{n_b})$  ▷ sequential, grouped, or FFT-tree
15:     $p_b \leftarrow m_b(y_b)$ 
16:     $\ell_b \leftarrow -\log p_b$ 
17:  end if
18: end for
19: return  $\mathcal{L} = B^{-1} \sum_{b=1}^B \ell_b, \{p_b\}_{b=1}^B$ 
```

151 **Two immediate consequences.** First, when the atomic support is Bernoulli, the convolutional
152 likelihood is exactly the coefficient of the polynomial $\prod_i ((1 - p_i) + p_i t)$ and therefore recovers
153 the standard count probability. Second, for any finite support, the likelihood uses the entire induced
154 aggregate distribution, not only the mean aggregate value. Thus two instance predictors with the
155 same expected sum can receive different likelihoods if they assign different probability mass to the
156 observed aggregate. This is the distinction exploited by the digit-sum and signed-count experiments.

157 4.2 Training Objective

158 For each minibatch, the instance model and task-specific contribution map first produce atomic
159 PMFs for every instance in every bag. The aggregate layer then deterministically maps these
160 atoms to an exact PMF over the observed bag label. This separation makes the aggregate objec-
161 tive backbone-agnostic: an MLP, CNN, transformer, or pretrained image encoder can be used as
162 long as it produces per-instance scores from which Φ can form the required finite-support PMFs. It
163 also makes the weak-supervision family modular: binary counts, signed counts, ordinal sums, and
164 histogram counts use different kernels but the same aggregate likelihood rule. The training loss is
165 the negative log probability assigned to the observed aggregate,

$$\mathcal{L}(\theta) = -\frac{1}{B} \sum_{b=1}^B \log (d_{b1} * \dots * d_{bn_b})(y_b). \quad (14)$$

166 Gradients are backpropagated through the convolutional aggregate layer into the instance network.
167 Instance labels are never used for training in the weak-supervision experiments; they are used only
168 for evaluation of hidden instance recovery. This is different from expected-value regression base-
169 lines, which regress the scalar expectation $\sum_i \mathbb{E}[z_i]$ to the aggregate label and therefore discard the
170 higher-order distributional information in the count or sum.

171 4.3 Special Cases

172 Table 1 summarizes the task-specific kernels; the aggregator itself is unchanged.

173 **Binary counts.** A target concept gives hidden Bernoulli indicators z_i , and the observed label is
174 $Y = \sum_i z_i$. The kernel $[1 - p_i, p_i]$ recovers the standard Bernoulli count likelihood, so this case is
175 included as a connection to prior count-loss work.

Table 1: Task kernels used by the same convolutional aggregator. The support order is the array order for the atomic PMF.

Task	Aggregate label	Atomic kernel	Support
Binary count	$Y = \sum_i z_i$	$[1 - p_i, p_i]$	$\{0, 1\}$
Signed count	$Y_{\pm} = \sum_i s_i z_i$	$\kappa_i^{(+)} = [0, 1 - p_i, p_i]; \kappa_i^{(-)} = [p_i, 1 - p_i, 0]$	$\{-1, 0, +1\}$
Integer sum	$Y_{\text{sum}} = \sum_i z_i$	$[p_{i0}, \dots, p_{i,K-1}]$	$\{0, \dots, K-1\}$
OVR histogram	$h_c = \sum_i z_{ic}$	$[1 - p_{ic}, p_{ic}]$ for each class c	$\{0, 1\}$ per class

Signed counts. With observed signs $s_i \in \{-1, +1\}$, target instances contribute s_i and non-target instances contribute 0, giving $Y_{\pm} = \sum_i s_i z_i$. The signed kernels in Table 1 expose cancellation because many latent assignments can yield the same signed sum.

Integer-valued sums. For ordinal labels $z_i \in \{0, \dots, K-1\}$, the atomic kernel is the full categorical vector and the observed label is $Y_{\text{sum}} = \sum_i z_i$. The likelihood therefore compares the observed sum to the full convolved PMF rather than only to the expected sum $\sum_i \mathbb{E}[z_i]$.

Histograms. For multiclass histograms, we use classwise one-vs-rest (OVR) Bernoulli count likelihoods: for each class c , the aggregator scores the observed count $h_c = \sum_i z_{ic}$, and the loss averages over classes. This is a scalable classwise likelihood, not a full joint multiclass histogram likelihood, and is closely related to the count-likelihood component of LLP-PVC.

4.4 Computational Backends

All of the above objectives are one-dimensional convolutions over integer aggregate support. If $m^{(i)}$ denotes the aggregate PMF after processing the first i atoms, then

$$m^{(i)}(s) = \sum_{r=a}^b m^{(i-1)}(s-r) d_i(r), \quad m^{(0)} = \delta_0. \quad (15)$$

This recurrence is the same probability law whether implemented as dynamic programming, shifted dense tensor accumulation, grouped convolution, or FFT polynomial multiplication. The choice of backend is therefore computational rather than statistical: the likelihood being optimized is unchanged, but the hardware profile differs. We use sequential convolution for moderate binary, signed, and integer-valued sum experiments; Shukla-style dynamic programming as a Bernoulli count baseline; grouped GPU convolution for signed Bernoulli groups; and a balanced FFT-tree backend for large one-vs-rest histogram PMFs. Further implementation details and backend complexity are given in Appendix A.

5 Experiments

5.1 Experimental Setup

We evaluate finite-support aggregation under five aggregate-label themes:

1. ordinal integer-valued sum labels;
2. signed count labels with cancellation;
3. Bernoulli count labels as a recovery of prior count loss;
4. class-count and histogram labels connecting to LLP;
5. runtime of exact aggregate backends.

Unless otherwise stated, MNIST-family experiments use LeNet-style instance networks, bags are sampled on the fly, and tables report means over three random seeds. Hidden instance labels are never used for training; they are used only for evaluation. Implementation details are summarized in Appendix A. Each subsection changes the observed aggregate label while keeping the same finite-support convolutional likelihood interface.

Table 2: Ordinal sum labels. Exact likelihood uses the full convolved PMF over sums, whereas MSE regresses only the expected sum. SVHN rows use ResNet-18 backbones; UltraMNIST reports best-checkpoint mean $\pm$ std. over five seeds.

Data	Bag/train	Method	Sum MAE	Instance acc.
MNIST	10/1000	exact sum	0.612	.982
MNIST	10/1000	MSE sum	2.535	.226
MNIST	10/5000	exact sum	0.277	.992
MNIST	10/5000	MSE sum	1.023	.746
MNIST	50/1000	exact sum	4.901	.693
MNIST	50/1000	MSE sum	6.274	.248
MNIST	50/5000	exact sum	2.229	.887
MNIST	50/5000	MSE sum	2.594	.817
SVHN	10/5000	exact sum, scratch	5.812	.229
SVHN	10/5000	MSE sum, scratch	5.328	.249
SVHN	10/5000	exact sum, pretr.	1.324	.955
SVHN	10/5000	MSE sum, pretr.	2.109	.826
UltraMNIST	3–5/800	exact sum	.341 $\pm$.038	.974 $\pm$.003

Metrics. Across the experimental tables, aggregate MAE columns report the mean absolute difference between the predicted aggregate and the observed aggregate over test bags. For expected-value baselines, the predicted aggregate is the scalar expectation trained by the baseline. For exact-likelihood rows, the prediction rule uses the aggregate PMF statistic specified by the corresponding table: count, integer-valued sum, signed count, or histogram count. For histograms, count MAE is the mean absolute class-count error averaged over classes and bags,

$$\frac{1}{N} \sum_B \frac{1}{C} \sum_{c=1}^C |\hat{h}_{Bc} - h_{Bc}|. \quad (16)$$

Rows labelled MSE denote MSE training objectives; MAE and accuracy columns are evaluation metrics computed after training. Columns labelled instance accuracy use held-out instance labels only for evaluation, never for training. Instance accuracy is never used for model selection unless explicitly stated.

5.2 Ordinal Sum Labels

Digit-sum supervision is an integer-valued aggregate label: the learner observes a single scalar while each hidden instance label has ten possible values. Here z_i is the original MNIST digit value in $\{0, \dots, 9\}$ and the bag label is $Y_{\text{sum}} = \sum_i z_i$. Exact finite-support likelihood uses the full induced distribution over possible sums, whereas an expected-value baseline regresses only $\sum_i \mathbb{E}[z_i]$ to the observed sum.

For $n = 10$ and 1000 training bags, the MSE baseline has substantially worse aggregate error and hidden digit recovery than the exact likelihood. With more training bags it improves, but the exact likelihood remains stronger in this table. For $n = 50$, the MSE baseline can recover some hidden digit structure while still making large aggregate errors, because small expectation errors accumulate across many digits. SVHN and UltraMNIST support the same conclusion beyond MNIST: pretraining makes natural RGB house-number digit sums usable, and exact likelihood remains substantially stronger than expected-sum regression even with the same pretrained ResNet-18 backbone. On SVHN with 5000 bags, the pretrained exact-likelihood model reaches 95.46% hidden digit accuracy and 65.00% exact sum accuracy, while the pretrained MSE baseline reaches 82.58% hidden digit accuracy and 22.58% exact sum accuracy. UltraMNIST gives a controlled patch-based benchmark with high hidden digit recovery.

5.3 Signed Aggregates with Cancellation

Signed counts test an aggregate that is not reducible to ordinary class proportions. We again use the digit-9 target indicator, but each instance carries a known random sign; target-digit images contribute +1 or −1, while all other digits contribute 0. Even with random signs, the observed signed sum is many-to-one: positive and negative contributions can cancel, so different latent assignments can

Table 3: Signed aggregates with random known signs. MSE regresses the expected signed sum $\sum_i s_i p_i$, while exact signed uses the full finite-support likelihood. MNIST rows report three-seed means; FashionMNIST rows report seed 0–1 means; CIFAR-10 rows are seed-0 pretrained ResNet-18 probes.

Data	Bag/train	Method	Signed MAE	Instance acc.
MNIST	10/1000	MSE signed	0.269	0.959
MNIST	10/1000	exact signed	0.095	0.990
MNIST	10/5000	MSE signed	0.047	0.995
MNIST	10/5000	exact signed	0.045	0.996
MNIST	50/1000	MSE signed	0.398	0.991
MNIST	50/1000	exact signed	0.243	0.994
MNIST	50/5000	MSE signed	0.203	0.995
MNIST	50/5000	exact signed	0.198	0.996
FashionMNIST	$n = 10$	MSE signed	0.121	0.987
FashionMNIST	$n = 10$	exact signed	0.113	0.988
FashionMNIST	$n = 50$	MSE signed	0.413	0.988
FashionMNIST	$n = 50$	exact signed	0.381	0.990
CIFAR-10	$n = 16$	exact signed, pretr.	0.343	0.974
CIFAR-10	$n = 64$	exact signed, pretr.	0.753	0.980

Table 4: Bernoulli count labels. Count likelihood recovers accurate hidden target detectors on MNIST target-count bags.

Data	Bag/train	Method	Count MAE	Instance acc.
MNIST	10/1000	exact count	0.115	0.988
MNIST	10/5000	exact count	0.090	0.991
MNIST	50/1000	exact count	0.686	0.987
MNIST	50/5000	exact count	0.460	0.991

yield the same bag label. The exact likelihood nevertheless yields lower signed-count error than a simple expected-signed-sum MSE baseline.

Both methods recover accurate target detectors with enough data, but the likelihood uses the full signed aggregate distribution rather than only $\sum_i s_i p_i$ and gives lower signed-count error. FashionMNIST shows the same pattern beyond MNIST: MSE is close in instance recovery, but exact likelihood is consistently better on the signed aggregate. The CIFAR-10 pretrained probes further show that signed finite-support aggregation can recover target-class detectors in a harder natural-image setting, reaching instance AUCs of 98.23% for $n = 16$ and 99.19% for $n = 64$. This setting is a direct example of why finite-support aggregation is not only a reformulation of ordinary non-negative counts: cancellation changes the inverse problem even when the hidden instance variable remains Bernoulli.

5.4 Recovery of Bernoulli Count Loss

Binary Count-MIL validates the connection to prior count-loss work. For MNIST-Bags, we make MNIST binary by choosing digit 9 as the positive class and treating digits 0–8 as negative. The learner observes the number of target-digit images in the bag; hidden instance labels are used only for evaluation. This case verifies that the convolutional formulation recovers the established Bernoulli count probability while sharing the same implementation pattern as the new aggregate types.

5.5 Class Histograms and LLP Connections

Histogram supervision is included mainly to connect the proposed finite-support view to recent LLP-PVC work. For MNIST histograms, the hidden category is the original digit class and the observed label is the vector of digit counts in the bag. We evaluate the classwise one-vs-rest count likelihood (OVR count), which corresponds to the Bernoulli count-likelihood component of LLP-PVC, rather

Table 5: Representative histogram-supervision results. Count MAE is mean absolute class-count error. OVR count denotes the classwise one-vs-rest count likelihood in Section 4. CIFAR-10 16/1000 MSE is final over seeds 0–1; CIFAR-10 64/1000 MSE is a running seed 0–1 mean around epoch 413 of 500.

Data	Bag/train	Backbone	Method	Count MAE	Instance acc.
MNIST	50/5000	LeNet	MSE	1.684	.116
MNIST	50/5000	LeNet	OVR count	.080	.992
CIFAR-10	16/1000	pretr. R18	MSE	.702	.668
CIFAR-10	16/1000	pretr. R18	OVR count	.260	.902
CIFAR-10	64/1000	pretr. R18	MSE	1.855	.272
CIFAR-10	64/1000	pretr. R18	OVR count	.678	.914
CIFAR-10	64/250	pretr. R18	OVR count	1.115	.803
CIFAR-10	64/250	pretr. R18	best by hist. MAE	1.013	.806

Table 6: Runtime for multiclass one-vs-rest count PMFs with batch size 256 and $C = 10$ classes. CPU DP is compared to GPU FFT-tree convolution.

Bag size	CPU DP (s)	GPU FFT-tree (s)	Max abs. error
64	8.246	0.00097	3.4×10^{-7}
128	17.230	0.00097	4.2×10^{-7}
256	36.785	0.00121	5.1×10^{-7}
512	54.040	0.00134	6.6×10^{-7}

than a full joint multiclass histogram likelihood. The comparison here is deliberately simple: MSE matches expected class counts, while OVR count scores the observed count with the exact classwise count PMF.

On the representative MNIST large-bag setting, OVR count likelihood improves over MSE proportion matching in count MAE and hidden digit accuracy; the omitted smaller-bag rows follow the same pattern. This supports the connection to class-count LLP objectives without making the histogram setting the main novelty claim. The CIFAR-10 rows are narrower: they show that OVR count likelihood is effective with ImageNet-pretrained ResNet-18 features and that it substantially outperforms a simple pretrained MSE histogram baseline at $n = 16$ and $n = 64$. However, the $n = 64$, train 1000 setting exposes four times as many image instances per epoch as $n = 16$, train 1000. The fixed-instance-budget setting, $n = 64$ with 250 train bags, reaches only about 80.3% final instance accuracy and 80.6% when selecting by observed histogram MAE. Thus the CIFAR-10 results should be read as evidence that count likelihood can use aggregate-labeled image volume effectively under a strong representation, not as evidence that larger bags are intrinsically superior.

5.6 Backend Runtime

Our FFT-tree backend computes batched multiclass one-vs-rest count PMFs on GPU and gives large practical speedups over CPU dynamic programming in large-bag benchmarks. This addresses a practical limitation of DP-based count losses: exact count probabilities need not be tied to a sequential CPU implementation. The benchmark isolates aggregate-PMF evaluation; it measures backend speed rather than end-to-end training time or representation quality.

6 Discussion

The experiments support a hierarchy of aggregate supervision. The largest gains occur when the observed aggregate is not well represented by only an expectation or a class proportion. Scalar expected-value matching is weak in low-data digit-sum settings, whereas exact aggregate likelihood can recover accurate instance predictors from the same scalar labels. With more training bags and small bags, expected-value baselines can become competitive, so the observed advantage is data efficiency and likelihood fidelity rather than universal dominance. The SVHN and UltraMNIST results strengthen the digit-sum evidence: the same integer-valued sum likelihood remains useful when the support is fixed at $S = 10$ but the image domain and backbone change. The signed FashionMNIST and CIFAR-10 probes make the same point for cancellation aggregates: the method

is not tied to MNIST digits or to nonnegative counts. Together with binary, ordinal, signed, and histogram benchmarks, these results support the view that FS-Conv is a finite-support aggregate-likelihood layer rather than a collection of task-specific losses.

Histogram labels are substantially richer, and standard LLP objectives can already be strong. We include the histogram setting mainly to connect the finite-support view to LLP-PVC-style one-vs-rest count likelihoods. Our main comparison shows that exact classwise count likelihood improves over MSE expected-count matching, while we do not claim a new full joint histogram likelihood or a reproduction of full LLP-PVC.

The same formulation is modular with respect to the instance model. Backbone choice determines the quality of the atomic PMFs, while the aggregate likelihood is determined by the weak label: binary counts, signed counts, integer-valued sums, and one-vs-rest histograms all use the same convolutional bag-prediction interface. This modularity is useful computationally as well as conceptually, because the exact aggregate PMF can be evaluated by CPU DP, direct convolution, grouped convolution, or GPU FFT-tree multiplication without changing the statistical objective.

Several limitations remain. Large bags can expose calibration issues even when instance ranking is accurate. Signed aggregates can be ambiguous because different latent assignments may yield the same observed sum. CIFAR-scale experiments require matching the backbone initialization and training protocol used by recent LLP work. Pretrained ResNet-18 one-vs-rest count likelihood is effective for $n = 16$ and $n = 64$, but the fixed-instance-budget comparison shows that the large-bag gain disappears when total sampled image exposure is matched. The CIFAR signed-count result is currently a single-seed probe, so it supports feasibility on natural images but should not be read as a final benchmark claim. These results show effective use of aggregate-labeled image volume under a strong representation, but they do not indicate intrinsic large-bag superiority or superiority to fully supervised learning at matched annotation budget.

References

- [1] Stuart Andrews, Ioannis Tsochantaridis, and Thomas Hofmann. Support vector machines for multiple-instance learning. In *Advances in Neural Information Processing Systems*, 2002.
- [2] Lorne Applebaum, Travis Dick, Claudio Gentile, Haim Kaplan, and Tomer Koren. Optimal learning from label proportions with general loss functions, 2025.
- [3] Gerda Bortsova, Florian Dubost, Silas Orting, Ioannis Katramados, Laurens Hogeweg, Laura Thomsen, Mathilde Wille, and Marleen de Bruijne. Deep learning from label proportions for emphysema quantification. *Medical Image Computing and Computer Assisted Intervention*, pages 768–776, 2018.
- [4] Robert Busa-Fekete, Heejin Choi, Travis Dick, Claudio Gentile, and Andres Munoz Medina. Easy learning from label proportions. In *Advances in Neural Information Processing Systems*, pages 14957–14968, 2023.
- [5] Robert Busa-Fekete, Travis Dick, Claudio Gentile, Haim Kaplan, Tomer Koren, and Uri Stemmer. Nearly optimal sample complexity for learning with label proportions, 2025.
- [6] Marc-Andre Carbonneau, Veronika Cheplygina, Eric Granger, and Ghyslaine Gagnon. Multiple instance learning: A survey of problem characteristics and applications. *Pattern Recognition*, 77:329–353, 2018.
- [7] James W. Cooley and John W. Tukey. An algorithm for the machine calculation of complex fourier series. *Mathematics of Computation*, 19(90):297–301, 1965.
- [8] Thomas G. Dietterich, Richard H. Lathrop, and Tomas Lozano-Perez. Solving the multiple instance problem with axis-parallel rectangles. *Artificial Intelligence*, 89(1–2):31–71, 1997.
- [9] Gabriel Dulac-Arnold, Neil Zeghidour, Marco Cuturi, Lucas Beyer, and Jean-Philippe Vert. Deep multi-class learning from label proportions, 2019.
- [10] Yen-Chang Hsu, Zhaoyang Lv, Joel Schlosser, Phillip Odom, and Zsolt Kira. Multi-class classification without multi-class labels. In *International Conference on Learning Representations*, 2019.

- [11] Maximilian Ilse, Jakub M. Tomczak, and Max Welling. Attention-based deep multiple instance learning. In *International Conference on Machine Learning*, 2018.
- [12] Jiabin Liu, Bo Wang, Zhiquan Qi, Yingjie Tian, and Yong Shi. Learning from label proportions with generative adversarial networks. In *Advances in Neural Information Processing Systems*, 2019.
- [13] Jiabin Liu, Bo Wang, Xin Shen, Zhiquan Qi, and Yingjie Tian. Two-stage training for learning from label proportions. In *International Joint Conference on Artificial Intelligence*, pages 2737–2743, 2021.
- [14] Tianhao Ma, Han Chen, Juncheng Hu, Yungang Zhu, and Ximing Li. Forming auxiliary high-confident instance-level loss to promote learning from label proportions. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 20592–20601, 2025.
- [15] Tianhao Ma, Wei Wang, Ximing Li, Gang Niu, and Masashi Sugiyama. Learning from label proportions via proportional value classification. In *International Conference on Learning Representations*, 2026.
- [16] Oded Maron and Tomas Lozano-Perez. A framework for multiple-instance learning. In *Advances in Neural Information Processing Systems*, 1998.
- [17] Novi Quadrianto, Alex J. Smola, Tiberio S. Caetano, and Quoc V. Le. Estimating labels from label proportions. *Journal of Machine Learning Research*, 10:2349–2374, 2009.
- [18] Julian Schroeter, Kirill A. Sidorov, and David Marshall. Weakly-supervised temporal localization via occurrence count learning. In *International Conference on Machine Learning*, pages 5649–5659, 2019.
- [19] Clayton Scott and Jianxin Zhang. Learning from label proportions: A mutual contamination framework. In *Advances in Neural Information Processing Systems*, volume 33, pages 22256–22267, 2020.
- [20] Vinay Shukla, Zhe Zeng, Kareem Ahmed, and Guy Van den Broeck. A unified approach to count-based weakly supervised learning. In *Advances in Neural Information Processing Systems*, volume 36, pages 38709–38722, 2023. URL https://proceedings.neurips.cc/paper_files/paper/2023/hash/79a0c8e7ae8e403e39341ea6b0ba4c21-Abstract-Conference.html.
- [21] Tao Sun, Dan Sheldon, and Brendan O’Connor. A probabilistic approach for learning with label proportions applied to the us presidential election. In *IEEE International Conference on Data Mining*, pages 445–454, 2017.
- [22] Kuen-Han Tsai and Hsuan-Tien Lin. Learning from label proportions with consistency regularization. In *Asian Conference on Machine Learning*, pages 513–528, 2020.
- [23] Xinggang Wang, Yongluan Yan, Peng Tang, and Wenyu Liu. Revisiting multiple instance neural networks. *Pattern Recognition*, 74:15–24, 2018.
- [24] Zixi Wei, Lei Feng, Bo Han, Tongliang Liu, Gang Niu, Xiaofeng Zhu, and Heng Tao Shen. A universal unbiased method for classification from aggregate observations. In *International Conference on Machine Learning*, pages 36804–36820, 2023.
- [25] Haoran Yang, Wanjing Zhang, and Wai Lam. A two-stage training framework with feature-label matching mechanism for learning from label proportions. In *Asian Conference on Machine Learning*, pages 1461–1476, 2021.
- [26] Felix X. Yu, Krzysztof Choromanski, Sanjiv Kumar, Tony Jebara, and Shih-Fu Chang. On learning from label proportions, 2014.
- [27] Jianxin Zhang, Yutong Wang, and Clayton Scott. Learning from label proportions by learning with label noise. In *Advances in Neural Information Processing Systems*, pages 26933–26942, 2022.

Table 7: Exact aggregate backends. n is bag size, S is atomic support size, $T = n(S - 1) + 1$ is aggregate PMF width, and C is number of one-vs-rest classes.

Backend	Setting	Work per bag	Use
DP	Bernoulli count	$O(nT)$	Shukla-style count baseline
Sequential convolution	finite support	$O(nST)$	main moderate-bag experiments
Grouped convolution	signed Bernoulli	$O(nT)$, groups	signed GPU implementation
FFT-tree	finite support / OVR	$O(CT \log T \log n)$	runtime, OVR histograms

A Implementation Details

All weak-supervision experiments sample bags on the fly rather than pregenerating fixed bag files. Instance labels are used only for evaluation metrics such as hidden-instance AUC or hidden digit accuracy. MNIST-family experiments use LeNet-style instance classifiers, Adam with learning rate 5×10^{-4} , weight decay 10^{-4} , and 50 epochs. The main MNIST-family configurations use seeds 0, 1, 2, test on 1000 sampled bags, and evaluate mean bag sizes 10 and 50 with standard deviations 2 and 10.

Backend implementation. Sequential finite-support convolution accumulates shifted copies of the current aggregate PMF and is used for the moderate-size binary, signed, and integer-valued sum experiments. Dynamic programming is used as a Shukla-style Bernoulli count baseline and computes the same binary count coefficients as the convolutional formulation. For grouped signed binary aggregates, each active group starts as a point mass at zero on support $[-K_{\max}, K_{\max}]$. For one signed instance, the grouped GPU backend applies a length-three channelwise kernel and padding one. Because PyTorch’s `conv1d` operator is cross-correlation rather than mathematical convolution, the implementation stores the kernel in the order

$$[p_i \mathbf{1}\{s_i = +1\}, 1 - p_i, p_i \mathbf{1}\{s_i = -1\}], \quad (17)$$

so the left kernel entry shifts probability mass toward larger aggregate values and the right entry shifts mass toward smaller values. This is only an implementation orientation; in mathematical support order $[-1, 0, +1]$ the signed kernels are the atoms described in Section 4.

For digit-sum, the atom width is the number of finite digit values, for example ten for MNIST digits. For histogram experiments, we avoid a full C -dimensional joint histogram over classes and instead compute C one-vs-rest one-dimensional count likelihoods in parallel. This keeps the aggregate support one-dimensional for each class and makes large bags practical on GPU. The FFT-tree backend is useful for this setting because the same pairwise polynomial products are batched over classes and bags. Padding or masked positions are represented by deterministic zero-contribution atoms, such as $[1, 0]$ for Bernoulli counts, $[0, 1, 0]$ for signed support $[-1, 0, +1]$, and $[1, 0, \dots, 0]$ for integer-valued sums.

Histogram and CIFAR settings. Histogram count-likelihood experiments treat each class count as a one-vs-rest Bernoulli count likelihood. The MNIST histogram experiments use batch size 192 for large-bag configurations. CIFAR-10 experiments use ResNet-18 backbones. The pretrained one-vs-rest count settings use ImageNet-pretrained ResNet-18 with random crop and horizontal flip augmentation, SGD with momentum 0.9, weight decay 10^{-4} , learning rate 5×10^{-4} , and a cosine learning-rate schedule for 500 epochs. The pretrained CIFAR-10 analysis is limited to the evaluated $n = 16$ and $n = 64$ one-vs-rest count settings.

NeurIPS Paper Checklist

The checklist is designed to encourage best practices for responsible machine learning research, addressing issues of reproducibility, transparency, research ethics, and societal impact. Do not remove the checklist: **The papers not including the checklist will be desk rejected.** The checklist should follow the references and follow the (optional) supplemental material. The checklist does NOT count towards the page limit.

Please read the checklist guidelines carefully for information on how to answer these questions. For each question in the checklist:

- You should answer [Yes], [No], or [N/A].
- [N/A] means either that the question is Not Applicable for that particular paper or the relevant information is Not Available.
- Please provide a short (1–2 sentence) justification right after your answer (even for [N/A]).

The checklist answers are an integral part of your paper submission. They are visible to the reviewers, area chairs, senior area chairs, and ethics reviewers. You will also be asked to include it (after eventual revisions) with the final version of your paper, and its final version will be published with the paper.

The reviewers of your paper will be asked to use the checklist as one of the factors in their evaluation. While [Yes] is generally preferable to [No], it is perfectly acceptable to answer [No] provided a proper justification is given (e.g., error bars are not reported because it would be too computationally expensive” or “we were unable to find the license for the dataset we used”). In general, answering [No] or [N/A] is not grounds for rejection. While the questions are phrased in a binary way, we acknowledge that the true answer is often more nuanced, so please just use your best judgment and write a justification to elaborate. All supporting evidence can appear either in the main paper or the supplemental material, provided in appendix. If you answer [Yes] to a question, in the justification please point to the section(s) where related material for the question can be found.

IMPORTANT, please:

- **Delete this instruction block, but keep the section heading “NeurIPS Paper Checklist”,**
- **Keep the checklist subsection headings, questions/answers and guidelines below.**
- **Do not modify the questions and only use the provided macros for your answers.**

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).

- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer **[N/A]** means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer **[N/A]** means that the paper does not include experiments.
- The authors should answer **[Yes]** if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer **[N/A]** means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.

- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines?>

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that the paper poses no such risks.

- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. Declaration of LLM usage

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy in the NeurIPS handbook for what should or should not be described.